



Reinforcement Learning and Related Topics



Goal: Present a brief overview of reinforcement learning and related topics along with resources to make it easier to dig deeper.



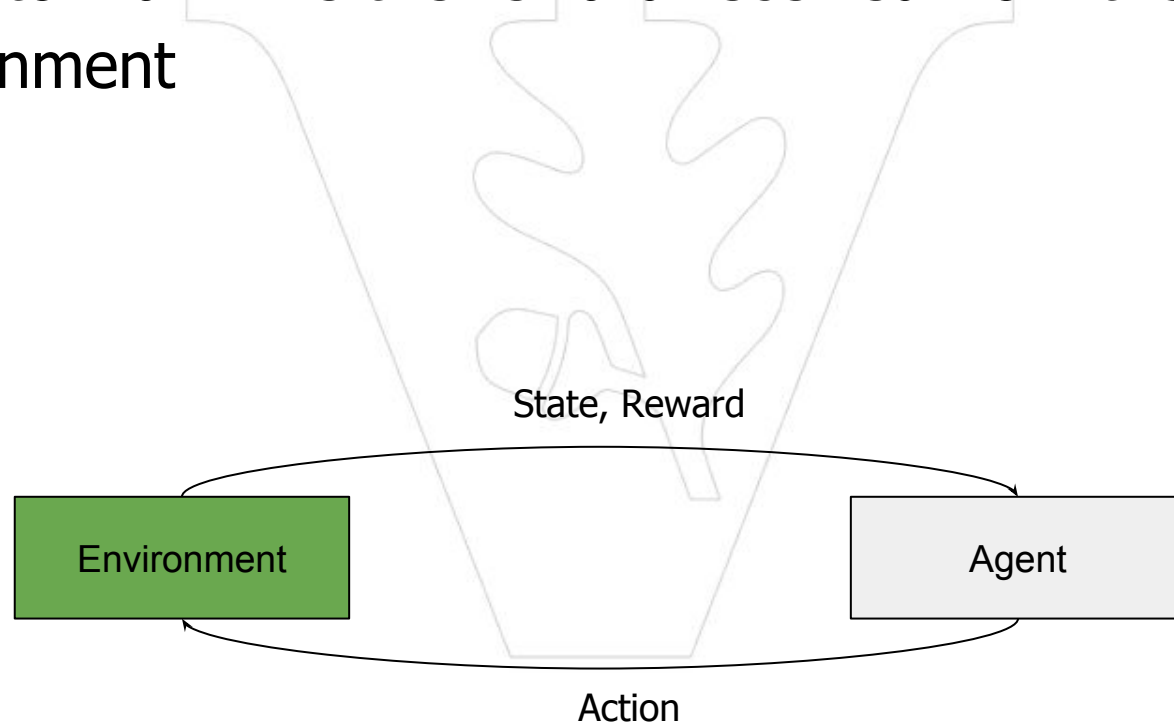
Reinforcement Learning



Basics



- There is an agent interacting in an environment and trying to maximize the reward received from the environment





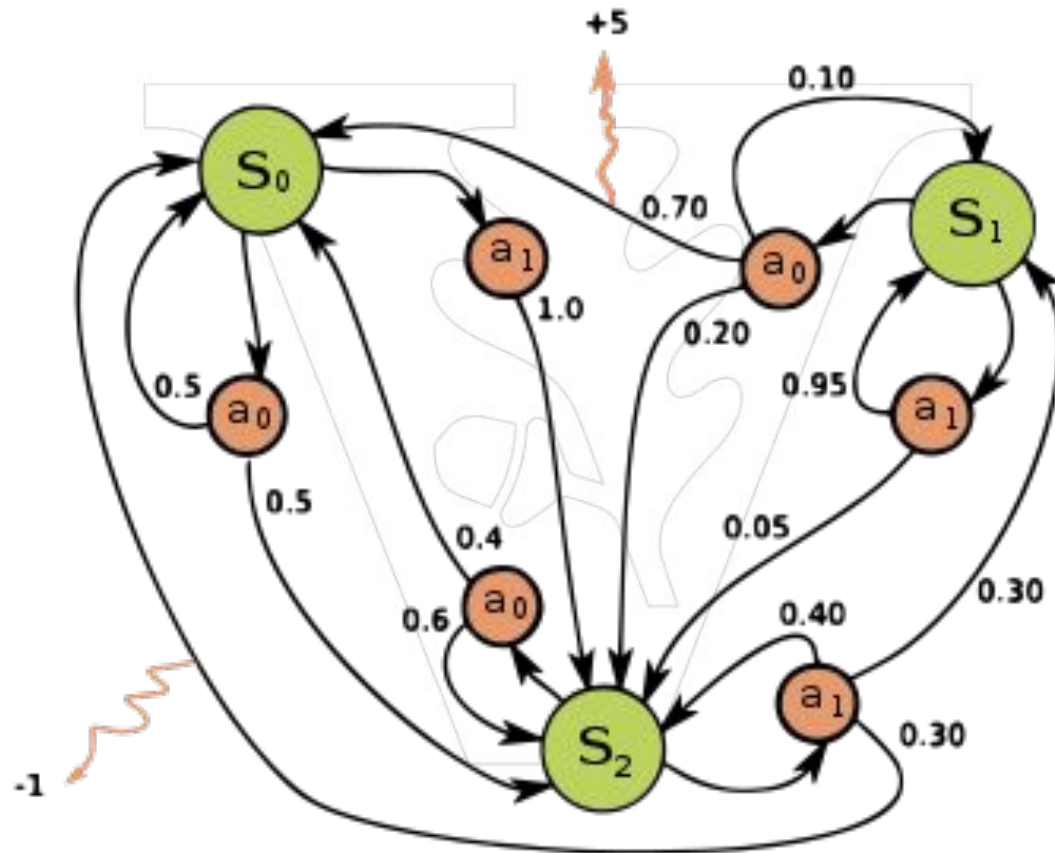
Terminology



- *State* is the numerical representation of the current state of the environment.
- An *episode* is one sequence of interactions (ie, trajectory) until termination. In games, this would mean playing until game over.
- *Reward* is a scalar returned after each interaction.
- *Return* is the cumulative reward over the course of a single episode
- An *agent* is the entity trying to maximize return.
- A *policy* defines the behavior for an agent.
- *State space* defines all possible values for the state.
- *Action space* defines all possible values for actions.



Markov Decision Process





Value and Q Functions



- Q function: Given a state and action, return the expected return.
- Value function: Given a state, return the expected reward from the state.
- [More Info...](#)

$$Q(s_t, a_t) = R$$

$$V(s_t) = R$$



A Few Algorithms






Before starting the example...



- OpenAI Gym has basically become the standard interface for RL environments.
- It makes environments interchangeable as they each define:
 - A state space
 - An action space
 - Reward function
 - Methods for applying actions and resetting the environment



CartPole in PyTorch





Related Topics



- What if interactions with the environment are expensive? Could we just learn the transition dynamics? (ie, replace the environment with a neural network)
 - Yep, this is *model-based reinforcement learning*
- **Deterministic vs stochastic environments.** If the environment is deterministic, you *could* memorize the best action sequence and simply replay it each time.
- What if the optimal policy is no longer optimal? This could mean that the environment is **non-stationary** making the problem much harder!
- **Exploration vs exploitation trade-off.** Some algorithms use approaches like epsilon-greedy exploration; others have exploration built into the algorithm.
- **On-policy vs off-policy.** Do we need data collected from our policy to improve (on-policy)? Or could we use data collected from other policies/agents/users (off-policy)?



More Related Topics



- What if actions in my environment do not impact subsequent decisions?
 - Then a **multi-armed bandit** is more suited for your task!
- What if I have human demonstrations available? Can I use them to help with the RL problem?
 - Yeah, you should be able to “warm start” the model using imitation learning initially then switch to reinforcement learning.
- What if I have a policy which is non-differentiable? Can I still optimize it?
 - There are optimization techniques such as CMAES which can be used to perform direct policy search.
- What if I want the policy to learn to be human-like?
 - You are probably looking for **imitation learning**



Demos and Resources



- [DQN Tutorial in PyTorch](#)
- [Spinning Up in Deep RL](#)
- [Bandit Algorithms Book](#)
- [RLlib \(framework with many implemented algorithms\) \(paper\)](#)
- [Imitation Learning using Case-based Reasoning in NetsBlox](#)